

16-350

Planning Techniques for Robotics

***Interleaving Planning and Execution:
Incremental Heuristic Search***

Maxim Likhachev

Robotics Institute

Carnegie Mellon University

Planning during Execution

- Planning is a repeated process!
 - partially-known environments
 - dynamic environments
 - imperfect execution of plans
 - imprecise localization
- Need to be able to re-plan fast!
- Several methodologies to achieve this:
 - anytime heuristic search: return the best plan possible within T msecs
 - **incremental heuristic search: speed up search by reusing previous efforts**
 - real-time heuristic search: plan few steps towards the goal and re-plan later

Planning during Execution

- Planning is a repeated process!
 - partially-known environments → *edgcost changes*
 - dynamic environments → *edgcost changes, goal changes*
 - imperfect execution of plans → *robot pose changes/deviates off the path*
 - imprecise localization → *robot pose changes/deviates off the path*
- Need to be able to re-plan fast!
- Several methodologies to achieve this:
 - anytime heuristic search: return the best plan possible within T msecs
 - **incremental heuristic search: speed up search by reusing previous efforts**
 - real-time heuristic search: plan few steps towards the goal and re-plan later

Only Goal Changes

Any ideas how to handle it?

Only Goal Changes


Any ideas how to handle it?

*Re-compute heuristics with respect to the **new** goal, and
continue searching until the **new** goal state is expanded*

Only Goal Changes

- Example on the board!

$$h(\text{cell } \langle x, y \rangle) = |x - x_{\text{goal}}| + |y - y_{\text{goal}}| \text{ (Manhattan Distance)}$$



Only Robot Pose Changes

Any ideas how to handle it?

Only Robot Pose Changes

Any ideas how to handle it?

Do the search backwards:

Then, the problem becomes “Only Goal Changes” that we know how to solve already


What if both Robot Pose and its Goal change?

Too bad!

Typically, you are better off re-planning from scratch then.

Changes to Edgecosts

- Two main reasons
 - Noisy perception (e.g., flickering obstacles, sensed position of obstacles is shifting, robot localization is noisy, etc.)
 - Partially-known environment




Changes to Edgecosts

*Typically, it is important to do clever filtering
to minimize flicker as much as possible without sacrificing safety*

- Two main reasons

- Noisy perception (e.g., flickering obstacles, sensed position of obstacles is shifting, robot localization is noisy, etc.)
- Partially-known environment



Changes to Edgecosts

*Typically, it is important to do clever filtering
to minimize flicker as much as possible without sacrificing safety*

- Two main reasons

- Noisy perception (e.g., flickering obstacles, sensed position of obstacles is shifting, robot localization is noisy, etc.)
- Partially-known environment

What should we assume about unknown space?

Changes to Edgecosts

- Two main reasons

Typically, it is important to do clever filtering to minimize flicker as much as possible without sacrificing safety


- Noisy perception (e.g., flickering obstacles, sensed position of obstacles is shifting, robot localization is noisy, etc.)
- Partially-known environment

What should we assume about unknown space?

***Freespace Assumption:** Assume that any “unknown” space is traversable until sensed otherwise!*

Changes to Edgecosts


- The robot doesn't initially know the status of the door



We ran an uninformed A^ search backwards
(that is, all g -values are costs to s_{goal})*

Changes to Edgcosts

- The robot doesn't initially know the status of the door





Why backwards?

We ran an uninformed A^ search backwards
(that is, all g -values are costs to s_{goal})*

Changes to Edgecosts


- The robot doesn't initially know the status of the door

4-connected grid



States with
changed
g-values

during execution, the robot found out that the door is closed

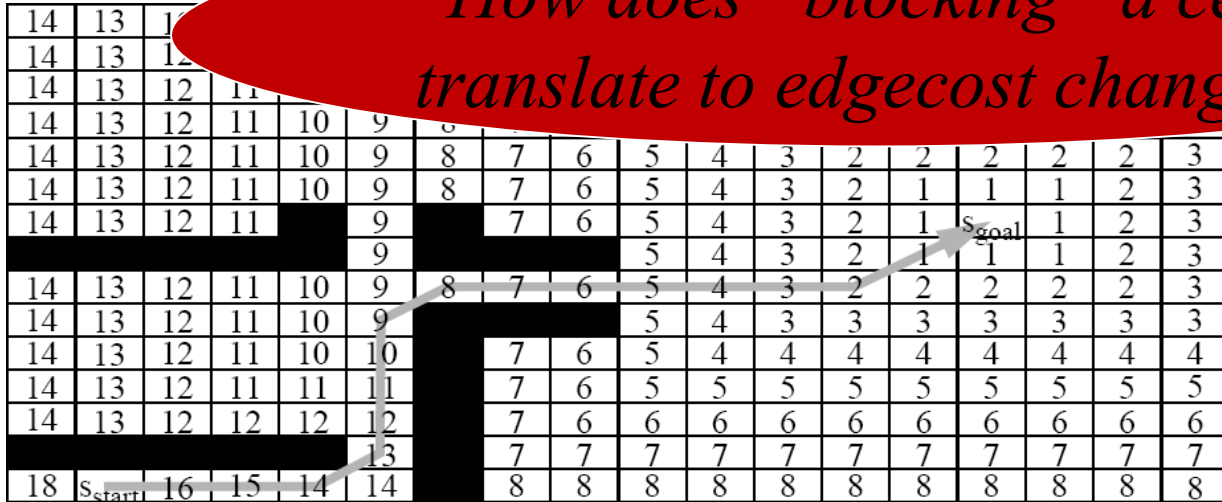



Changes to Edgecosts

- The robot doesn't initially know the status of the door


How does “blocking” a cell translate to edgecost changes?

4-connected grid



States with changed g-values



during execution, the robot found out that the door is closed



Changes to Edgecosts

- The robot doesn't initially know the status of the door


4-connected grid



States with
changed
g-values

during execution, the robot for


Can we reuse these g-values from one search to another? – incremental A*




Incremental Heuristic Search

- D^*/D^* Lite: Incremental Heuristic Search Algorithms


*initial search by backwards A^**




initial search by D^ Lite*



*second search by backwards A^**




second search by D^ Lite*



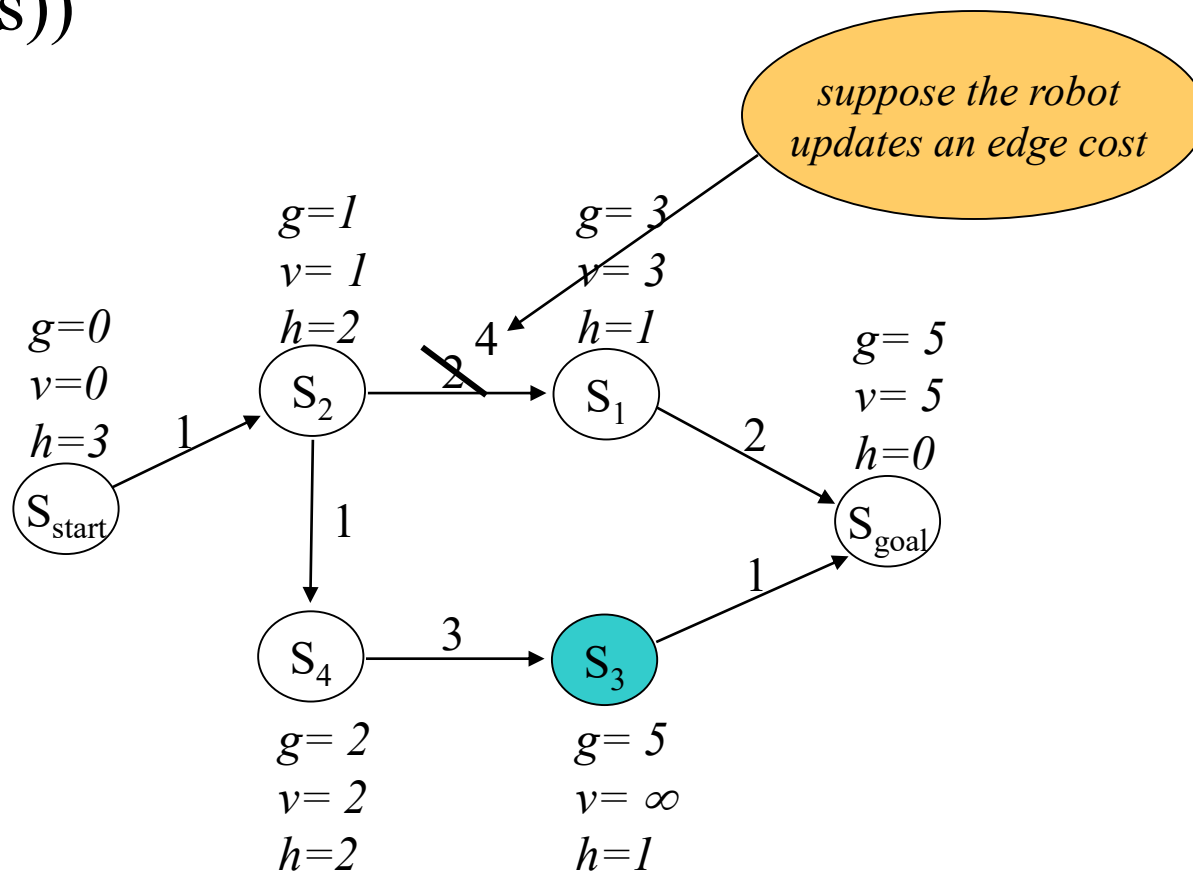
A* with Reuse of State Values

- So far, ComputePathwithReuse() could only deal with states whose $v(s) \geq g(s)$ (overconsistent or consistent)
- Edge cost increases may introduce underconsistent states ($v(s) < g(s)$)



A* with Reuse of State Values

- So far, ComputePathwithReuse() could only deal with states whose $v(s) \geq g(s)$ (overconsistent or consistent)
- Edge cost increases may introduce underconsistent states ($v(s) < g(s)$)




A* with Reuse of State Values

- Edge cost increases may introduce underconsistent states ($v(s) < g(s)$)

ComputePathwithReuse invariant:
 $g(s') = \min_{s'' \in \text{pred}(s')} v(s'') + c(s'', s')$

need to update $g(s_1)$




A* with Reuse of State Values

- Edge cost increases may introduce underconsistent states ($v(s) < g(s)$)

ComputePathwithReuse invariant:
 $g(s') = \min_{s'' \in \text{pred}(s')} v(s'') + c(s'', s')$

need to update $g(s_l)$


$v(s_l) < g(s_l)$



A* with Reuse of State Values

- Edge cost increases may introduce underconsistent states ($v(s) < g(s)$)
- Fix these by setting $v(s) = \infty$


ComputePathwithReuse invariant:
 $g(s') = \min_{s'' \in \text{pred}(s')} v(s'') + c(s'', s')$



A* with Reuse of State Values

- Edge cost increases may introduce underconsistent states ($v(s) < g(s)$)
- Fix these by setting $v(s) = \infty$
- Makes s overconsistent or consistent $v(s) \geq g(s)$

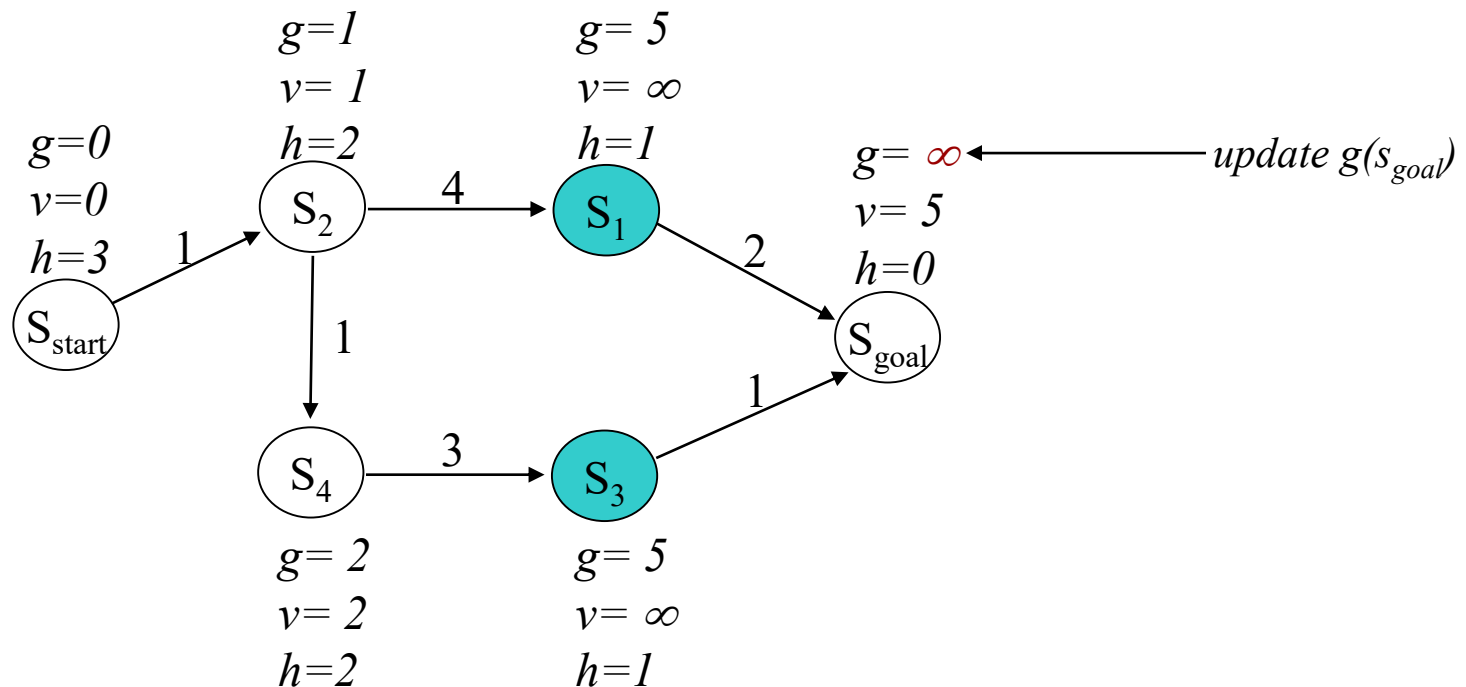
ComputePathwithReuse invariant:
 $g(s') = \min_{s'' \in \text{pred}(s')} v(s'') + c(s'', s')$



A* with Reuse of State Values

- Edge cost increases may introduce underconsistent states ($v(s) < g(s)$)
- Fix these by setting $v(s) = \infty$
- Makes s overconsistent or consistent $v(s) \geq g(s)$
- Propagate the changes

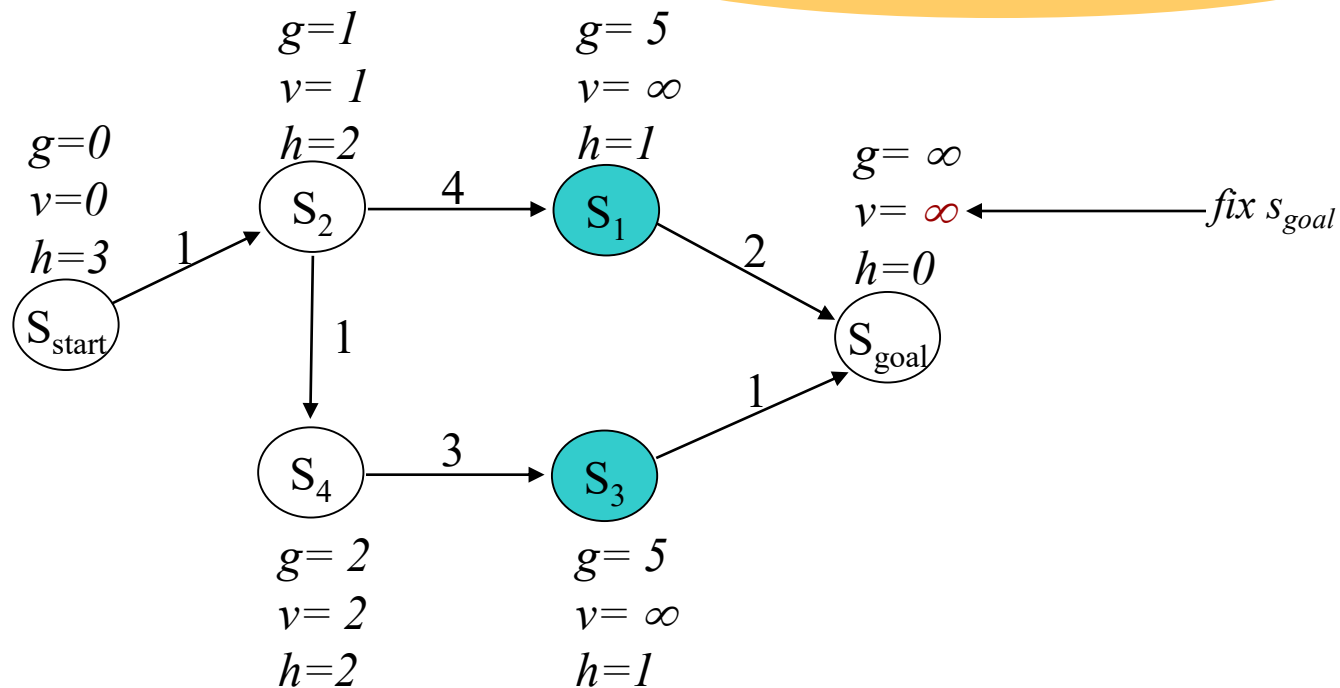
ComputePathwithReuse invariant:
 $g(s') = \min_{s'' \in \text{pred}(s')} v(s'') + c(s'', s')$



A* with Reuse of State Values

- Edge cost increases may introduce underconsistent states ($v(s) < g(s)$)
- Fix these by setting $v(s) = \infty$
- Makes s overconsistent or consistent $v(s) \geq g(s)$
- Propagate the changes

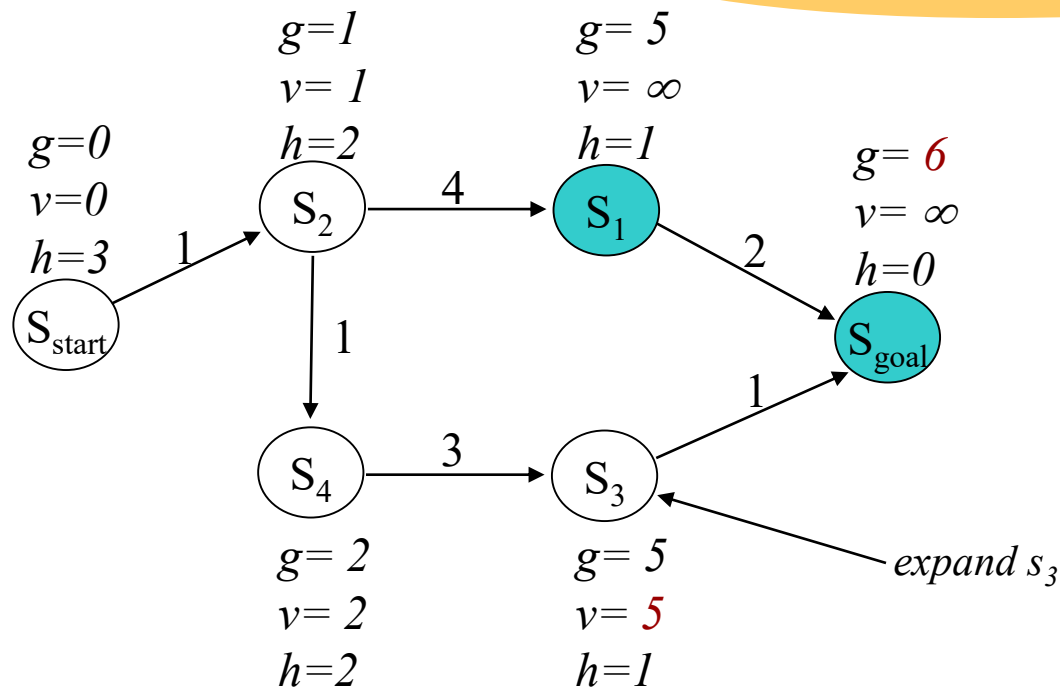
ComputePathwithReuse invariant:
 $g(s') = \min_{s'' \in \text{pred}(s')} v(s'') + c(s'', s')$



A* with Reuse of State Values

- Edge cost increases may introduce underconsistent states ($v(s) < g(s)$)
- Fix these by setting $v(s) = \infty$
- Makes s overconsistent or consistent $v(s) \geq g(s)$
- Propagate the changes

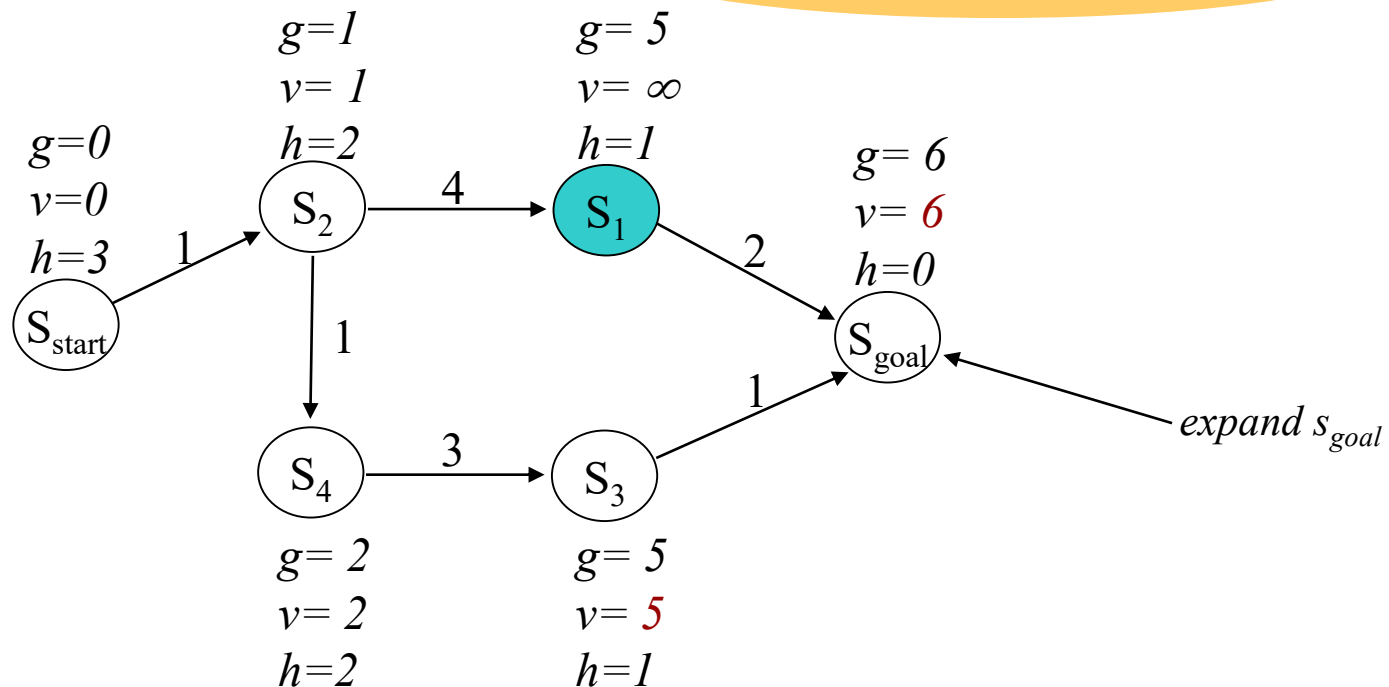
ComputePathwithReuse invariant:
 $g(s') = \min_{s'' \in \text{pred}(s')} v(s'') + c(s'', s')$



A* with Reuse of State Values

- Edge cost increases may introduce underconsistent states ($v(s) < g(s)$)
- Fix these by setting $v(s) = \infty$
- Makes s overconsistent or consistent $v(s) \geq g(s)$
- Propagate the changes

ComputePathwithReuse invariant:
 $g(s') = \min_{s'' \in \text{pred}(s')} v(s'') + c(s'', s')$

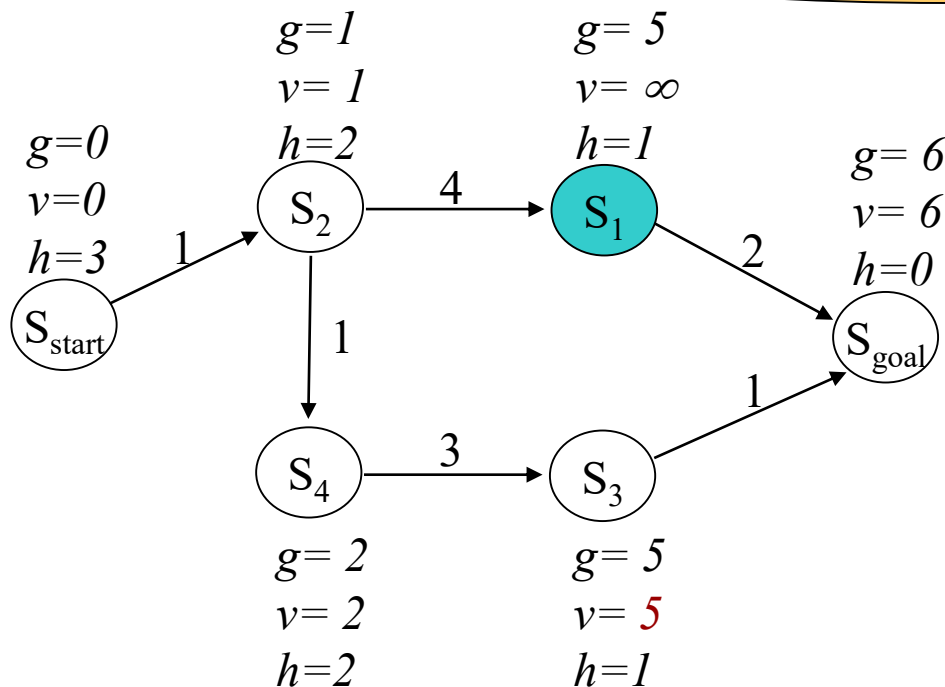


A* with Reuse of State Values

- Edge cost increases may introduce underconsistent states ($v(s) < g(s)$)
- Fix these by setting $v(s) = \infty$
- Makes s overconsistent or consistent
- Propagate the changes

*after ComputePathwithReuse terminates:
all g-values of states are equal to final A* g-values*

*we can backtrack an optimal path
(start at s_{goal} , proceed to pred that minimizes $g+c$)*

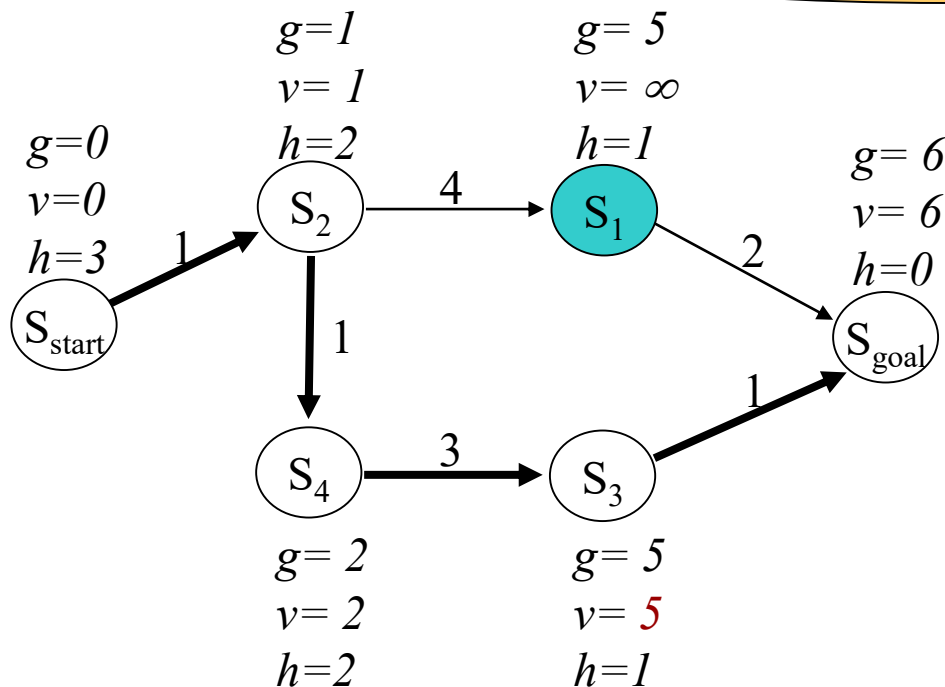


A* with Reuse of State Values

- Edge cost increases may introduce underconsistent states ($v(s) < g(s)$)
- Fix these by setting $v(s) = \infty$
- Makes s overconsistent or consistent
- Propagate the changes

*after ComputePathwithReuse terminates:
all g-values of states are equal to final A* g-values*

*we can backtrack an optimal path
(start at s_{goal} , proceed to pred that minimizes $g+c$)*



D* Lite

- Optimal re-planning algorithm
- Simpler and with nicer theoretical properties version of D*

until goal is reached

 ComputePathwithReuse(); *//modified to fix underconsistent states*

 publish optimal path;

 follow the path until map is updated with new sensor information;

 update the corresponding edge costs;

 set s_{start} to the current state of the agent;

Anytime Incremental Heuristic Search

- Anytime D*:
 - decrease ε and update edge costs at the same time
 - re-compute a path by reusing previous state-values

set ε to large value;

until goal is reached

 ComputePathwithReuse(); *//modified to fix underconsistent states*

 publish ε -suboptimal path;

 follow the path until map is updated with new sensor information;

 update the corresponding edge costs;

 set s_{start} to the current state of the agent;

 if significant changes were observed

 increase ε or replan from scratch;

 else

 decrease ε ;

What for?

Summary

- Changes to Start Only or Goal Only can be handled easily by continuing the search until “goal of the search” is expanded
- Freespace Assumption: assume that unknown is traversable until detected otherwise
- D*/D* Lite: Incremental versions of heuristic search (don't need to know the exact algorithm)